

# Distributed Data Systems

---

DIANE WOODBRIDGE, PH.D



UNIVERSITY OF SAN FRANCISCO  
CHANGE THE WORLD FROM HERE

# Announcement

## Homework

- Individual HW 3 - Feb 13 Midnight (Cannot be extended)
- Group HW2 - Feb 25th
  - See the instructions and videos
  - Make sure to watch Replica/Sharding Video by Feb 26th

## Quiz 1

- Multiple Choice

### Question 9

1 / 1 pts Regraged Score: 1 / 1 pts

! This question has been regraded.

When you execute the following commands, the first line returns 0. After executing the second line, what is the output of the third line?

```
db.class.find().count() // 1) Returns 0  
db.class.update({"code": "msds694"}, {"title": "distributed computing"}, {"upsert": true}) // 2)  
db.class.find().count() // 3)
```



0



1



2

None of the above

Additional Comments:



# Content

---

## MongoDB

- Aggregation Operations (aggregation pipeline)
- Index



# Content

---

## MongoDB

- **Aggregation Operations (aggregation pipeline)**
- Index



borough: "Manhattan"

grades:

[?date: —  
grade: A  
Score: 10],

?date: —  
grade: A  
Score: 83,

?date: —  
grade: B  
Score: 73 ]

borough: Brooklyn

grades:

[?date: —

grade: A  
Score: [1 ?]

borough:  
Manhattan

grades:

[?date: —,  
grade: A,  
Score: 83]

:

what's the avg score per grade  
for businesses in "Manhattan"?

# MongoDB

## Aggregation Pipeline

- Perform several stages of operations.
  - Each stage performs an operation on the input documents, and its outputs are passed to the next stage
  - An aggregation pipeline can return results for groups of documents.
  - Syntax
    - `db.collection_name.aggregate([{transform_operator_1 : criteria_1}, {transform_operator_2 : criteria_2}, ...])`
    - `transform_operators` performs operations including filtering, transformation, grouping, sorting using `$match`, `$project`, `$group`, `$unwind`, `$sort`, `$lookup` etc.



<https://docs.mongodb.com/manual/aggregation/>

<https://docs.mongodb.com/manual/reference/operator/aggregation-pipeline/>



# MongoDB

## Aggregation Pipeline

- Perform several stages of operations.
  - Each stage performs an operation on the input documents, and its outputs are passed to the next stage
  - An aggregation pipeline can return results for groups of documents.
  - `db.collection_name.aggregate([{$transform_operator_1 : criteria_1}, {$transform_operator_2 : criteria_2}, ...])`

Transform Operators	Description	Example
\$match	Filters the documents to allow only matching documents to pass the the next pipeline stage.	<code>db.collection.aggregate({\$match : {"size":{\$gt : 0}}})</code>
\$project	Reshapes each document in the stream, such as by adding new fields or removing existing fields.	<code>db.collection.aggregate({\$match : {"size":{\$gt : 0}}}, {\$project: {"size": false}})</code>
\$group	Groups input documents by a specified identifier expression.	<code>db.collection.aggregate({\$group:{_id :"date", count: {\$sum:1}}})</code>
\$unwind	Deconstructs an array field from the input documents to output a document for each element.	<code>db.collection.aggregate({\$unwind : "grade"})</code>
\$sort	Reorders documents by a specified sort key.	<code>db.collection.aggregate({\$sort:{date:-1}})</code>
\$lookup	Performs a left outer join to another collection in the same database.	<code>db.world_bank_project.aggregate({\$lookup: {from:"country_code", localField:"countrycode", foreignField:"code", as:"country_phone_info"}})</code>

# MongoDB

---

## Aggregation Pipeline

- Perform several stages of operations.
  - Each stage performs an operation on the input documents, and its outputs are passed to the next stage
  - An aggregation pipeline can return results for groups of documents.
  - Syntax
    - `db.collection_name.aggregate([{transform_operator_1 : criteria_1}, {transform_operator_2 : criteria_2}, ...])`
  - Field Path (**\$<field>**) in criteria
    - Aggregation pipelines use field path to access fields in the input documents.
      - "`$<field>`" is equivalent to "`$$CURRENT.<field>`" where the CURRENT is the root of the current object.
      - To access the value of the field, prefix the variable name with double dollar signs ( `$$` )

<https://docs.mongodb.com/manual/meta/aggregation-quick-reference/#field-paths>

<https://www.mongodb.com/docs/manual/reference/aggregation-variables/>

# MongoDB

---

## Aggregation Pipeline

- Pipeline transform operators
  - **\$match** : Filter document with criteria.
  - Format : { \$match: { <query> } }
  - The query syntax is identical to the read operation query syntax.

# Example 1

---

Using aggregate(), return documents where "id" is "P130164"



# Example 1

Using aggregate(), return documents where "id" is "P130164"

```
105 db.world_bank_project.aggregate({$match:{"id":"P130164"}})
```

Find × Aggregate × Aggregate × Find × Find × Aggregate × Text × Aggregate × Aggregate × Aggregate × Aggregate × Aggregate × »99+

|← ← → | 50 JSON View ⌂ ⚙

Documents 1 to 1

```
1 {
2   "_id" : ObjectId("52b213b38594d8a2be17c788"),
3   "approvalfy" : "2014",
4   "board_approval_month" : "October",
5   "boardapprovaldate" : "2013-10-29T00:00:00Z",
6   "borrower" : "THE GOVERNMENT OF INDIA",
7   "closingdate" : "2018-12-31T00:00:00Z",
8   "country_namecode" : "Republic of India!$!IN",
9   "countrycode" : "IN",
10  "countryname" : "Republic of India",
11  "countryshortname" : "India",
12  "docty" : "Project Appraisal Document,Integrated Safeguards Data Sheet,Procurement Plan,Procurement F
13  "
```



# MongoDB

---

## Aggregation Pipeline

- Pipeline transform operators
  - **\$project** : Extract field. Rename the projected field.
  - Project expressions
    - **Specify the inclusion of fields, the addition of new fields, and the resetting the values of existing fields.**
    - Syntax:{ \$project : { <specification(s)> } }
      - **specification**
        - {‘field’ : ‘true’/‘false’}
        - {‘new\_field’ : {\$operator: expression}}}

New fields can be included by various expressions including operations arithmetic, array, boolean, string expression operators and accumulators.

Ex. {\$project : {new\_field : {\$log: [number, base]}}} }

<https://docs.mongodb.com/manual/reference/operator/aggregation-pipeline/>

<https://docs.mongodb.com/manual/reference/operator/aggregation/project/>

<https://docs.mongodb.com/manual/reference/operator/aggregation/>

# MongoDB

## Aggregation Pipeline

- Operators which can be used to construct expression in transform operations.

Operators	Type	Description	Syntax
\$add	Arithmetic	Adds numbers to return the sum, or adds numbers and a date to return a new date.	{ \$add: [ "\$<field1>", "\$<field2>" ] }
\$eq/gt/lt	Comparison	Returns true if the values are equivalent/greater or less than the <value>.	{ \$gt: [ "\$<field>", <value> ] }
\$filter	Array	Selects a subset of an array to return based on the specified condition.	{ \$filter: { input:"\$<array_field>" , as: "<new-varname>", cond: <expression> } }
\$size	Array	Counts and returns the total number of items in an array.	{ \$size: "\$<array_field>" }
\$cond	Conditional	A ternary operator that evaluates one expression, and depending on the result, returns the value of one of the other two expressions.	{ \$cond: [ <boolean-expression>, <true-case>, <false-case> ] }
\$concat	String	Concatenates strings.	{ \$concat: [ "\$<field1>", "\$<field2>", ... ] }
\$dateDiff	Date	Returns the difference between two dates. Unit : year, month, day, hour, minute, etc.	{ \$dateDiff: { startDate: "\$<field1>", endDate: "\$<field2>", unit: "day" } }
\$convert	Type	Converts a value to a specified type.	{ \$convert : {input : "\$<field>", to: "data_type"} }

# Example 2

In world\_bank\_project, return "id" and a non-empty "human\_development\_project" array which represents "mjtheme\_namecode" only containing documents which name is "Human development".

**Expected Output**

```
{
  "id" : "P129828",
  "human_development_project" : [
    {
      "name" : "Human development",
      "code" : "8"
    }
  ]
}

{
  "id" : "P132616",
  "human_development_project" : [
    {
      "name" : "Human development",
      "code" : "8"
    }
  ]
}
```



# Example 2

In world\_bank\_project, return "id" and a non-empty "human\_development\_project" array which represents "mjtheme\_namecode" only containing documents which name is "Human development".

```
db.world_bank_project.aggregate(  
    {$project: {"_id": false, "id": true,  
               "human_development_project" :  
                   {$filter : {input: "$mjtheme_namecode",  
                               as : "mjtheme_namecode_a",  
                               cond: {$eq:[ "$$mjtheme_namecode_a.name",  
                                           "Human development"]}}},  
    })
```



```
{  
  "id" : "P144832",  
  "human_development_project" : [  
  ]  
}  
  
{  
  "id" : "P144665",  
  "human_development_project" : [  
  ]  
}
```



# Example 2

In world\_bank\_project, return "id" and a non-empty "human\_development\_project" array which represents "mjtheme\_namecode" only containing documents which name is "Human development".

```
db.world_bank_project.aggregate(  
    {$project: {"_id": false, "id": true,  
               "human_development_project" :  
                   {$filter : {input: "$mjtheme_namecode",  
                               as : "mjtheme_namecode_a",  
                               cond: {$eq:[ "$$mjtheme_namecode_a.name",  
                                           "Human development"]}}},  
        },  
    {$project: {"id":true, "human_development_project":true,  
               "size" : {$size : "$human_development_project"} } }  
)
```



```
{  
  "_id" : "P144832",  
  "human_development_project" : [  
  ],  
  "size" : NumberInt(0)  
}  
  
{  
  "_id" : "P144665",  
  "human_development_project" : [  
  ],  
  "size" : NumberInt(0)  
}
```



# Example 2

In world\_bank\_project, return "id" and a non-empty "human\_development\_project" array which represents "mjtheme\_namecode" only containing documents which name is "Human development".

```
db.world_bank_project.aggregate(  
    {$project: {"_id": false, "id": true,  
               "human_development_project" :  
                   {$filter : {input: "$mjtheme_namecode",  
                               as : "mjtheme_namecode_a",  
                               cond: {$eq:[ "$$mjtheme_namecode_a.name",  
                                           "Human development"]}}},  
        },  
    {$project: {"id":true, "human_development_project":true,  
               "size" : {$size : "$human_development_project"}},  
        {$match : {"size":{$gt : 0}}},  
        {$project: {"size": false}}  
    )
```



The screenshot shows the MongoDB Compass interface with the "Aggregate Query" tab selected. The query is displayed above the results. Below the results, there are buttons for "Pin Result", "Query Code", and "Explain Query". The results show a single document with an ID and a non-empty "human\_development\_project" array containing one element.

```
{  
  "id" : "P129828",  
  "human_development_project" : [  
    {  
      "name" : "Human development",  
      "code" : "8"  
    }  
  ]  
}
```

# MongoDB

---

## Aggregation Pipeline

- Pipeline operators
  - **\$group** : Group documents based on certain fields and combine their values.
  - Create/contain an \_id field which contains the distinct group by key.
  - Syntax
    - { \$group: { \_id: "\$<field>",  
                  <new\_field\_name>: { <accumulator1> : <expression1> }, ... } }
  - Accumulator operators can be used together. (also use with \$<field path>)
  - Ex. \$sum, \$avg, \$min, \$max, \$first, \$last, etc.

# Example 3

---

For each countrycode, return the number of documents in the collection in a format of `{"_id" : "code", "count" : value}`

## Expected Output

```
{  
    "_id" : "GT",  
    "count" : 2.0  
}  
  
{  
    "_id" : "CR",  
    "count" : 1.0  
}
```



# Example 3

For each countrycode, return the number of documents in the collection in a format of {"\_id": "code", "count" : value}

```
db.world_bank_project.aggregate({$group:{_id :"$countrycode", count:{$sum:1}}})
```

The screenshot shows the MongoDB Compass interface with the results of an aggregate query. The results are displayed in JSON format, grouped by country code and counting the number of documents.

```
{  
  "_id" : "GT",  
  "count" : 2.0  
},  
,  
{  
  "_id" : "CR",  
  "count" : 1.0  
},  
,  
{  
  "_id" : "BO",  
  "count" : 2.0  
}
```



# MongoDB

---

## Aggregation Pipeline

- Pipeline operators
  - **\$unwind** : Return each of an array element into a separate document.
  - Syntax
    - { \$unwind: '\$<field\_path>' }
    - **To specify a field path, prefix the field name with a dollar sign \$ and enclose in quotes.**

# Example 4

Return theme\_namecode as an \_id and the number of project has the corresponding theme in the entire collection.

## Expected Output

```
{  
    "_id" : {  
        "name" : "Judicial and other dispute resolution mechanisms",  
        "code" : "32"  
    },  
    "count" : 2.0  
}  
{  
    "_id" : {  
        "name" : "Other communicable diseases",  
        "code" : "64"  
    },  
    "count" : 3.0  
}
```



# Example 4

Return theme\_namecode as an \_id and the number of project has the corresponding theme in the entire collection.

```
db.world_bank_project.aggregate( { $unwind : "$theme_namecode" ,  
                                { $group : { _id : "$theme_namecode" ,  
                                             count:{$sum:1}} } ,  
                                )  
  
// Example 3  
[[{"Output": "Aggregate Query (line 38)", "X": 1}, {"Output": "Documents 1 to 50", "X": 2}, {"Output": "Pin Result", "X": 3}, {"Output": "Query Code", "X": 4}, {"Output": "Explain Query", "X": 5}], [{"Output": "Aggregate Query (line 38)", "X": 1}, {"Output": "Documents 1 to 50", "X": 2}, {"Output": "Pin Result", "X": 3}, {"Output": "Query Code", "X": 4}, {"Output": "Explain Query", "X": 5}],  
  
{  
  "_id" : {  
    "name" : "Judicial and other dispute resolution mechanisms" ,  
    "code" : "32"  
  } ,  
  "count" : 2.0  
}
```



# MongoDB

---

## Aggregation Pipeline

- **\$sort** : Collect all document and properly sort them.
- Syntax
  - { \$sort: { <field1>: <sort order>, <field2>: <sort order> ... } }
  - sort order : 1 (ascending), -1 (descending)

# Example 5

In world\_bank\_project, return the ‘approvalfy’ in double as ‘approvalfy\_num’ with ‘borrower’, and ‘id’ ordered by approvalfy in descending order.

## Expected Output

```
{  
    "borrower" : "GOVERNMENT OF TUNISIA",  
    "id" : "P144674",  
    "approvalfy_num" : 2015.0  
}  
  
{  
    "borrower" : "GOVERNMENT OF JORDAN",  
    "id" : "P144832",  
    "approvalfy_num" : 2014.0  
}
```

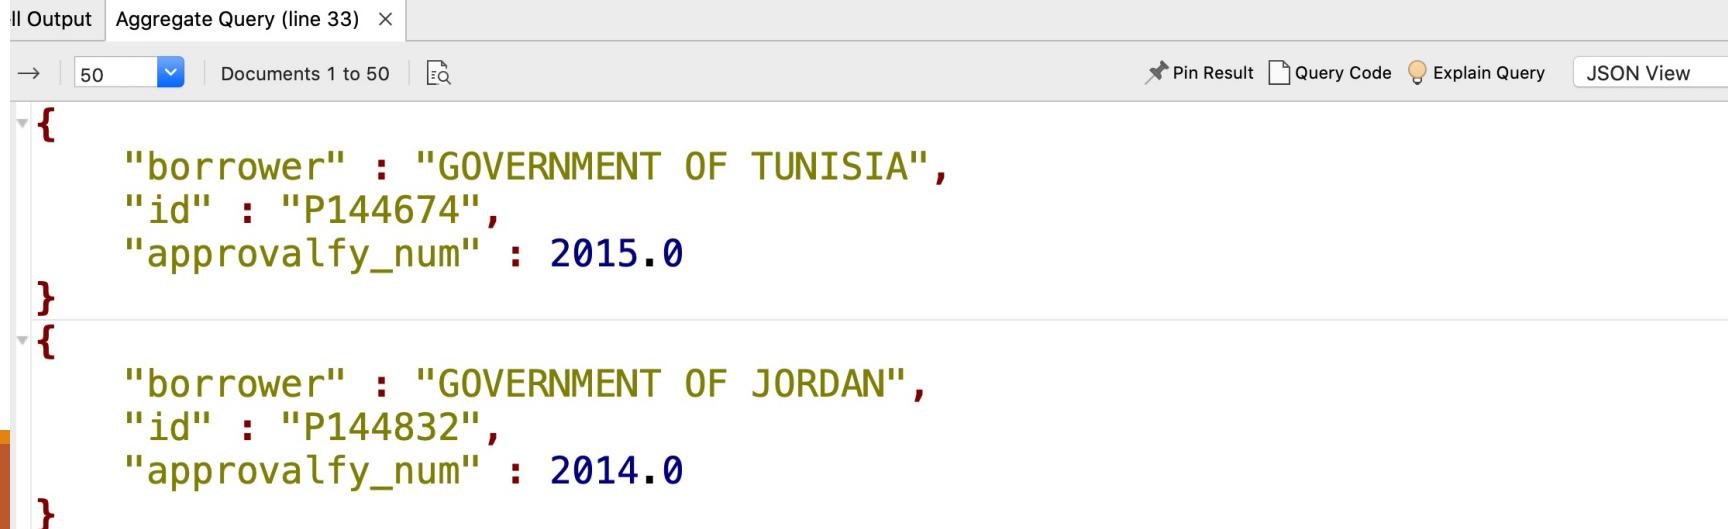
Operators	Type	Description	Syntax
\$convert	Type	Converts a value to a specified type.	{\$convert : {input : "\$<field>", to: "data_type"}}



# Example 5

In world\_bank\_project, return the ‘approvalfy’ in double as ‘approvalfy\_num’ with ‘borrower’, and ‘id’ ordered by approvalfy in descending order.

```
db.world_bank_project.aggregate({$project:  
    {"id" : true,  
     "borrower" : true,  
     "_id" : false,  
     "approvalfy_num" : {$convert:  
         {input:"$approvalfy", to:"double"}  
     }  
},  
    {$sort:{"approvalfy num": -1}})
```



The screenshot shows the MongoDB Compass interface with the results of the aggregate query. The results are displayed in a JSON array format, showing two documents. Each document contains the fields 'borrower', 'id', and 'approvalfy\_num'. The 'approvalfy\_num' field is a double value representing the original string 'approvalfy' converted from string to double.

```
{  
  "borrower" : "GOVERNMENT OF TUNISIA",  
  "id" : "P144674",  
  "approvalfy_num" : 2015.0  
}  
  
{  
  "borrower" : "GOVERNMENT OF JORDAN",  
  "id" : "P144832",  
  "approvalfy_num" : 2014.0  
}
```

# Example 6 - Extra

In world\_bank\_project, 1) find a document where totalcommamt is greater than 75000000 and 2) find the unique names of projectdocs' DocType and the number of projectdocs belonging to the sorted by DocType.

## Expected Output

```
{  
    "_id" : {  
        "DocType" : "AGR"  
    },  
    "count" : 27.0  
}  
  
{  
    "_id" : {  
        "DocType" : "AGRS"  
    },  
    "count" : 1.0  
}
```



# Example 6 - Extra

In world\_bank\_project, 1) find a document where totalcommamt is greater than 75000000 and 2) find the unique names of projectdocs' DocType and the number of projectdocs belonging to the sorted by DocType.

```
db.world_bank_project.aggregate({$match:{"totalcommamt": {$gte : 75000000}},  
                                {$project:{"projectdocs":1}},  
                                {$unwind:"$projectdocs"},  
                                {$group:{"_id": {"DocType":  
                                         "$projectdocs.DocType"},  
                                         "count": {$sum: 1}}},  
                                {$sort:{"_id.DocType":1}})
```



The screenshot shows the MongoDB Compass interface with the results of the aggregate query. The results pane displays the following JSON document:

```
{  
  "_id" : {  
    "DocType" : "AGR"  
  },  
  "count" : 27.0  
}
```



# MongoDB

---

## Aggregation Pipeline

- **\$lookup** : Performs a left outer join to a collection.  
It adds a new array field whose elements are the matching documents from the "joined" collection.
- Syntax
  - `db.<collection_1>.aggregate( { $lookup:{ localField: "<field_1>",  
from: "<collection_2>",  
foreignField: "<field_2>",  
as: "<output_field>"}} )`

# Example 7

---

Create a collection called country\_code using country\_codes.json.

For all the projects whose "countrycode" is "CN", create a field called country\_phone\_info which includes data from country\_code (includling "name", "dial\_code", "code", and "\_id").

```
        "url" : "http://www.worldbank.org/projects/P127033/china-renewab",
        "country_phone_info" : [
            {
                "_id" : ObjectId("61fbc97af5e71beac2ee0ce2"),
                "name" : "China",
                "dial_code" : "+86",
                "code" : "CN"
            }
        ]
    ]
```

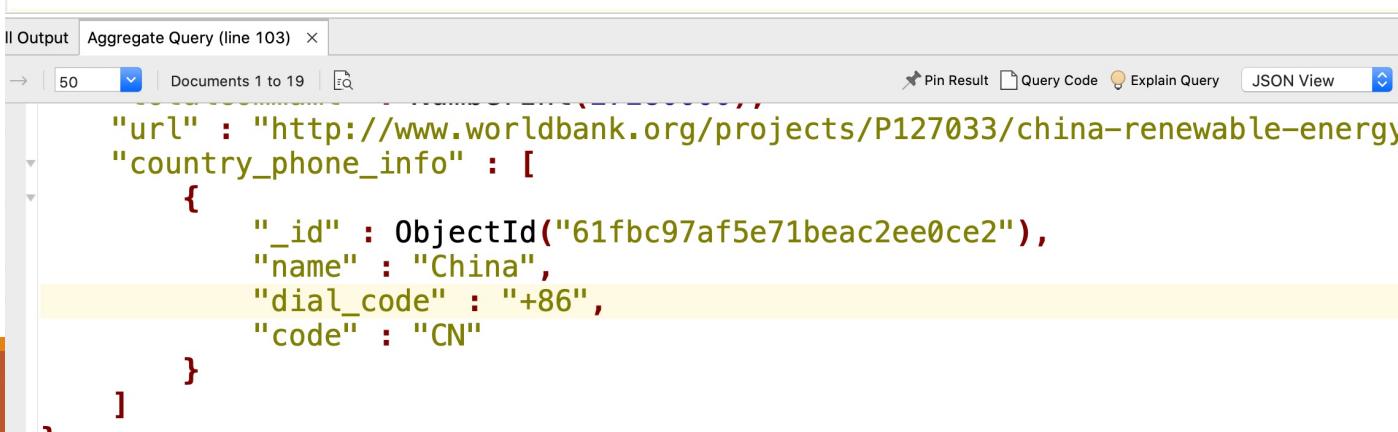


# Example 7

Create a collection called country\_code using country\_codes.json.

For all the projects whose "countrycode" is "CN", create a field called country\_phone\_info which includes data from country\_code (includling "name", "dial\_code", "code", and "\_id").

```
db.world_bank_project.aggregate([
    {$match:{"countrycode":"CN"}},
    {$lookup:{  
        from:"country_code",  
        localField:"countrycode",  
        foreignField:"code",  
        as:"country_phone_info"}  
    }]  
)
```



```
"url" : "http://www.worldbank.org/projects/P127033/china-renewable-energy",
"country_phone_info" : [
    {
        "_id" : ObjectId("61fbc97af5e71beac2ee0ce2"),
        "name" : "China",
        "dial_code" : "+86",
        "code" : "CN"
    }
]
```

# Content

---

## MongoDB

- Aggregation Operations (aggregation pipeline)
- **Index**



# Index

---

## INDEX

ABC, 164, 321n  
academic journals, 262, 280–82  
Adobe eBook Reader, 148–53  
advertising, 36, 45–46, 127, 145–46, 167–68, 321n  
Africa, medications for HIV patients in, 257–61  
Agee, Michael, 223–24, 225  
agricultural patents, 313n  
Aibo robotic dog, 153–55, 156, 157, 160  
AIDS medications, 257–60  
air traffic, land ownership vs., 1–3  
Akerlof, George, 232  
Alben, Alex, 100–104, 105, 198–99, 295, 317n  
alcohol prohibition, 200  
*Alice's Adventures in Wonderland* (Carroll), 152–53

Anello, Douglas, 60  
animated cartoons, 21–24  
antiretroviral drugs, 257–61  
Apple Corporation, 203, 264, 302  
architecture, constraint effected through, 122, 123, 124, 318n  
archive.org, 112  
*see also* Internet Archive  
archives, digital, 108–15, 173, 222, 226–27  
Aristotle, 150  
Armstrong, Edwin Howard, 3–6, 184, 196  
Arrow, Kenneth, 232  
art, underground, 186  
artists:  
    publicity rights on images of, 317n  
    recording industry payments to, 52,  
        58–59, 74, 195, 196–97, 199, 301,  
        329n–30n

- Organize data by a given field. → Optimize queries and enhance query performance.
- Things to consider.
  - Which fields to index?
  - Index scheme : balanced search tree (B-tree), hash, etc.



# Index

---

## Mongo DB

- Without indexes, MongoDB must perform a *collection scan*, i.e. scan every document in a collection, to select those documents that match the query statement.
- Index : Optimize query performance.
  - **Creating Index**
  - Syntax
    - `createIndex({“field” : value })`
    - **value**
      - Range-based indexes - Direction : 1 (Ascending), -1 (Descending)
        - MongoDB indexes use a B-tree data structure, by default.
      - Hashed indexes – “hashed”
      - Also supports other index types including text index, 2d index, etc.

<https://docs.mongodb.com/manual/indexes/>



# Index

---

## Mongo DB

- Without indexes, MongoDB must perform a *collection scan*, i.e. scan every document in a collection, to select those documents that match the query statement.
- Index : Optimize query performance.
  - **Describing** the existing indexes on the collection.
  - Syntax
    - `getIndexes()`
    - Check which indexes exist on a collection.
    - Default : `_id` (Cannot be deleted)

<https://docs.mongodb.com/manual/indexes/>



# Index

---

## Mongo DB

- Without indexes, MongoDB must perform a *collection scan*, i.e. scan every document in a collection, to select those documents that match the query statement.
- Index : Optimize query performance.
  - **Dropping or removing** the specified index from a collection.
  - Syntax
    - `dropIndex({“field” : direction})`

<https://docs.mongodb.com/manual/indexes/>



# Example 8

---

For world\_bank\_project,

- 1) Find a document where "countrycode" is "AO" and see its executionStats using  
`cursor.explain("executionStats")`.
- 2) Create a hashed-index on "countrycode".
- 3) Find a document where "countrycode" is "AO" and see its executionStats using  
`cursor.explain("executionStats")`.
- 4) Drop the index.



When poll is active, respond at **pollev.com/msds**

 Text **MSDS** to **37607** once to join

# Without an index, how many documents were examined for .find()?

500 (All)

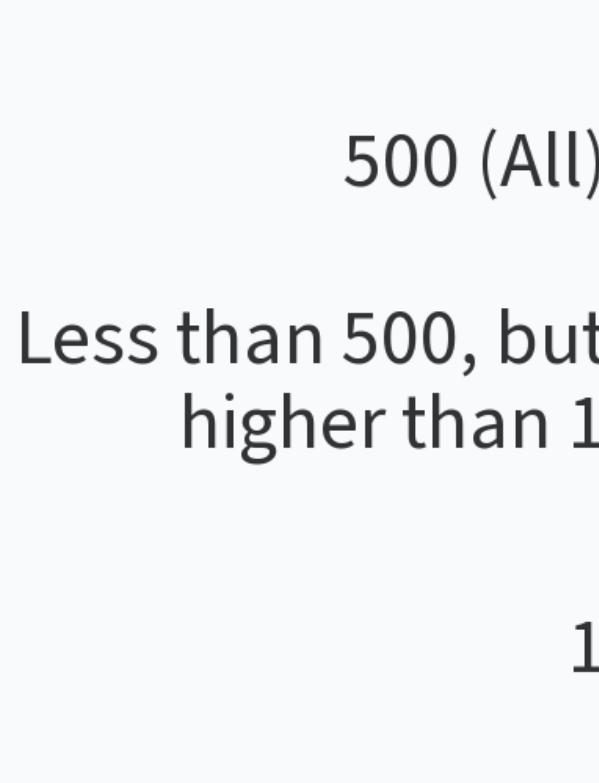
Less than 500, but higher than 1

1

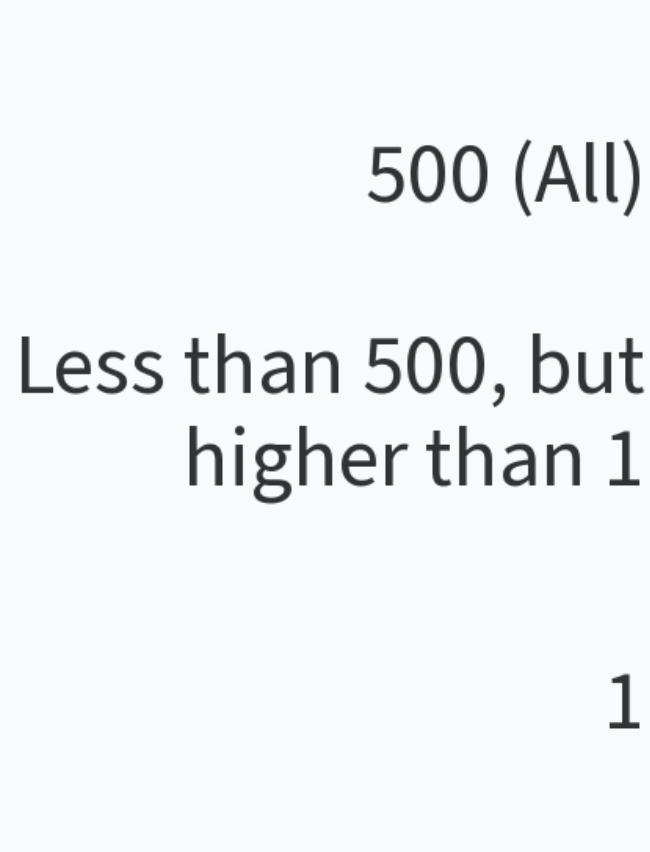
When poll is active, respond at **pollev.com/msds**

Text **MSDS** to **37607** once to join

# Without an index, how many documents were examined for .find()?



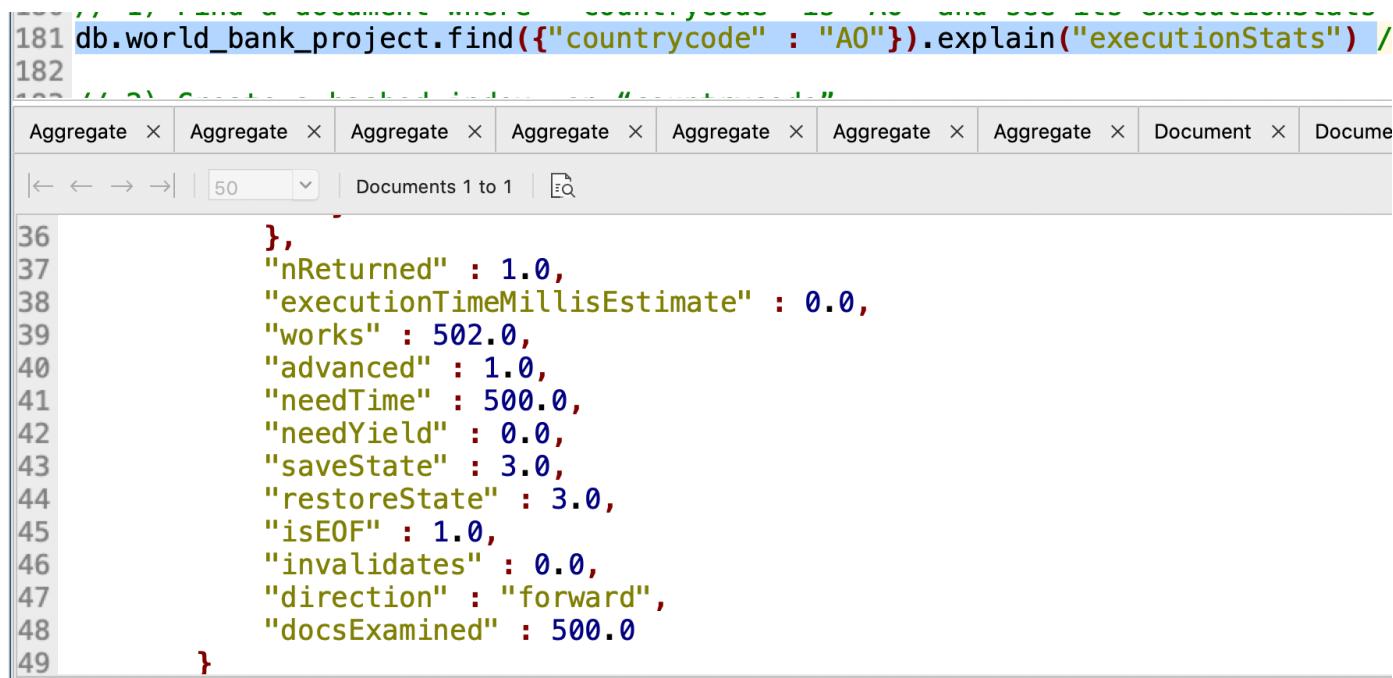
# Without an index, how many documents were examined for .find()?



# Example 8

For world\_bank\_project,

- 1) Find a document where "countrycode" is "AO" and see its executionStats using `cursor.explain("executionStats")`.



The screenshot shows a MongoDB shell session. The command entered is:

```
181 db.world_bank_project.find({"countrycode" : "AO"}).explain("executionStats") /
```

The results pane displays the execution statistics for the query. The execution stats object includes the following fields:

```
36     },
37     "nReturned" : 1.0,
38     "executionTimeMillisEstimate" : 0.0,
39     "works" : 502.0,
40     "advanced" : 1.0,
41     "needTime" : 500.0,
42     "needYield" : 0.0,
43     "saveState" : 3.0,
44     "restoreState" : 3.0,
45     "isEOF" : 1.0,
46     "invalidates" : 0.0,
47     "direction" : "forward",
48     "docsExamined" : 500.0
49 }
```



# Example 8

---

For world\_bank\_project,

- 2) Create a hashed-index on "countrycode".
- 3) Find a document where "countrycode" is "AO" and see its executionStats using `cursor.explain("executionStats")`.

```
183 // 2) Create a hashed-index on "countrycode".
189 // 4) Drop the index
190 db.world_bank_project.dropIndex({"countrycode": "hashed"})
```

Aggregate	Aggregate	Aggregate	Aggregate	Document	Document	Document
x	x	x	x	x	x	x

← ← → → | 50 | Documents 1 to 1 | ⚙

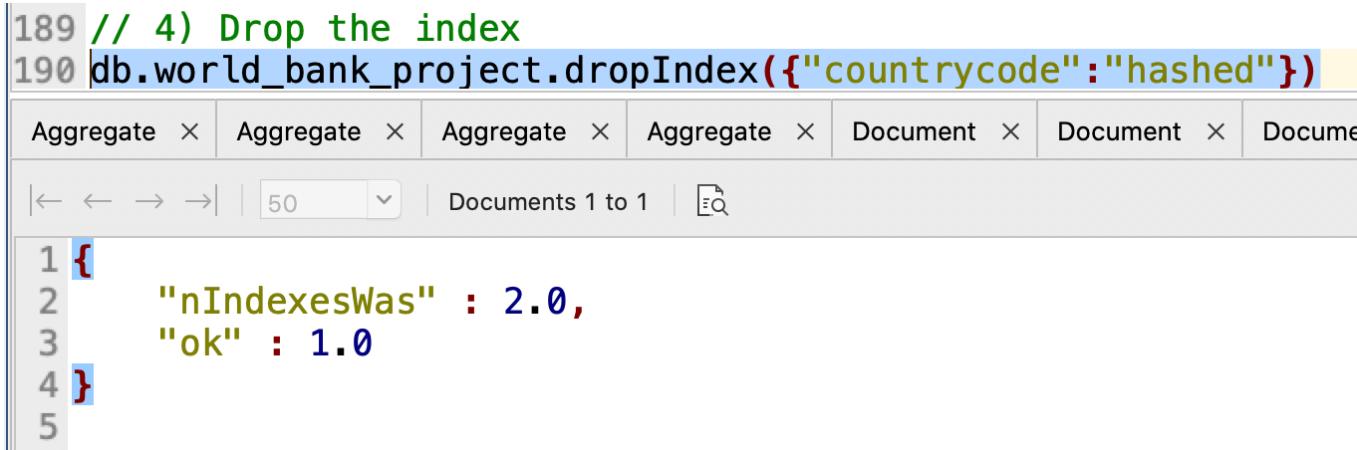
```
1 {
2   "nIndexesWas" : 2.0,
3   "ok" : 1.0
4 }
5
46   "totalDocsExamined" : 1.0,
47   "executionStages" : {
48     "stage" : "FETCH",
49     "filter" : {
50       "countrycode" : {
51         "$eq" : "AO"
52       }
53     }
54 }
```

# Example 8

---

For world\_bank\_project,

- 4) Drop the index.



The screenshot shows a MongoDB shell interface with the following code:

```
189 // 4) Drop the index
190 db.world_bank_project.dropIndex({"countrycode": "hashed"})
```

The command at line 190 is highlighted in blue. Below the code, the results of the command are displayed:

```
1 {
2   "nIndexesWas" : 2.0,
3   "ok" : 1.0
4 }
5
```

The results show that 2 indexes were present before the drop operation, and the operation was successful (ok: 1.0).



# Reference

---

1. MongoDB Online Documentation : <https://docs.mongodb.com/manual/tutorial/query-documents/>

